

# Introduction to Deep Neural Networks

Máster Universitario en Ciencia de Datos - Métodos Avanzados en Aprendizaje Automático

Carlos María Alaíz Gudín

Escuela Politécnica Superior  
Universidad Autónoma de Madrid

Academic Year 2021/22



Universidad Autónoma  
de Madrid



# Contents

① Introduction

② Innovations of Deep Learning



# Introduction



- ① MLPs (NNs with only one-hidden layer) were the state-of-the-art models during the 80s and 90s.
  - ② Due to the **universal approximation property** of the MLPs, deeper networks were not considered.
  - ③ With the apparition of Kernel Methods in the 90s, their use decreased considerably.
- 
- ④ They resurged again with the **Deep Learning** paradigm introduced by Hinton in 2006.

### Definition (Deep Learning)

**Deep Learning** (DL) is a type of machine learning based on artificial neural networks in which multiple layers of processing are used to extract progressively higher level features from data.



# MLPs with Several Hidden Layers: Limitations



- ① A main limitation of the MLPs with multiple hidden layers was the **vanishing gradient** problem.
  - The gradient tends to get smaller as it is propagated during the backward phase.
  - As a result, only the last layers are really trained, while the initial ones are kept almost unchanged.
  - It was easier to train an MLP with many hidden units in a single hidden layer than with many layers.

---

- ② Another difficulty is just the computational cost of training a large neural network.



Notebook

Vanishing Gradient



# MLPs with Several Hidden Layers: Diagnosis

- Hinton summarized why MLPs with several networks used to not work:
  - ① *Our labeled datasets were thousands of times too small.*
    - ⇒ Larger datasets.
  - ② *Our computers were millions of times too slow.*
    - ⇒ More computational power.
  - ③ *We initialized the weights in a stupid way.*
    - ⇒ Clever initialization.
  - ④ *We used the wrong type of non-linearity.*
    - ⇒ New activation functions.



# Innovations of Deep Learning





# Big Data



- The **Industry 4.0** and the **Digital Transformation** implied a revolution in the management of data.
  - The interest of the companies and institutions changed in two phases:
    - ① Collecting data and applying new technologies.
    - ② Trying to gather information and extracting value from the collected data.
- 
- This transformation resulted in the availability of a huge amount of heterogeneous data.
    - **Big data** paradigm.
- 
- The machine learning models can (and have to) be trained with much more data than before.



# Computational Power


- The computational power of CPUs has increased consistently, influenced among others by **Moore's law**.
  - The number of threads available per CPU has also raised.
- 
- A key factor in the development of DL is the usage of Graphics Processing Units (GPUs), which can handle hundreds of threads.
  - This allows for a huge degree of parallelization in matrix calculus.
- 
- Some companies (e.g. Google) has developed specific DL hardware as the Tensor Processing Units (TPUs).



CPU.



GPU.



TPU.



# Initialization



- Deep NNs can be properly trained if the weights are correctly initialized.
  - If they are too small, the gradient will vanish.
  - If they are too large, the learning can be very slow.
- There are several heuristics to initialize the weights effectively.

**Xavier Initialization** The weights are initialized using a Gaussian with zero mean and variance  $\frac{1}{d_{\ell-1}}$ , where  $d_{\ell-1}$  is the number of input units to each layer.

**Uniform Initialization** The weights are initialized using a uniform distribution around zero with bounds  $\pm\sqrt{2/(d_{\ell+1} + d_{\ell})}$  (the constant 2 depends on the activation function).

**Transfer Learning** The weights of a successfully trained model used in a similar problem are used as initial weights.



Notebook


Weight Initialization



# Activation Functions (I)

**Rectified Linear (ReLU)**  $g(\mathbf{w}^\top \mathbf{x} + b) = \max\{0, \mathbf{w}^\top \mathbf{x} + b\}$ .


- Sparse, the gradient does not vanish.
- Continuous but non-differentiable at 0.



## Activation Functions (II)

**Softplus**  $g(\mathbf{w}^\top \mathbf{x} + b) = \ln(1 + e^{\mathbf{w}^\top \mathbf{x} + b})$ .

- Smooth version of the ReLU.
- Continuous and differentiable, although non-sparse.



Notebook

Activation Functions



# Avoiding Over-Fitting: Data Augmentation



- With the increment in the flexibility of NNs a problem arises: the risk of **over-fitting**.
- 
- A large amount of data prevents from over-fitting.
    - Not always available, it depends on the problem.
- 
- A first solution is to generate new data (**data augmentation**).
  - This process is not trivial, the generated data has to be relevant for the problem.
  - Different approaches:
    - Perturbing with noise.
    - Fitting the distribution of the original data.
    - Using expert knowledge about the variations in real-life (particularly useful with images).





Notebook

Data Augmentation



# Avoiding Over-Fitting: Transfer Learning



- The model can be pre-trained in a large dataset, and then adapted to the problem at hand.
- This approach is known as **transfer learning**.

## Transfer Learning

- 1 Take a model successfully trained over a larger dataset.
  - The complete model, or only a part of it (usually the feature extraction).
- 2 Add the necessary layers for adapting it to the problem at hand.
- 3 Train the new layers.
- 4 Train all the layers with a smaller learning rate (**fine tuning**).



Notebook

Transfer Learning



# Avoiding Over-Fitting: Dropout



- A typical approach to regularize a DNN is the **dropout**.
- 
- 1 During training, a certain percentage  $r$  of the inputs to a hidden layer are set to 0, while the remaining inputs are increased as  $\frac{1}{1-r}$  to compensate the scale.
    - The network “learns” to distribute the information processing, not relying on single units.
  - 2 During the prediction of new data all the units are considered as usual.



Notebook

Dropout



# Specialized Architectures



- Another advance of the DL models is the use of **specialized architectures** designed for specific problems.
- Some examples are:
  - ① Autoencoders.
  - ② Convolutional Neural Networks.
  - ③ Recurrent Neural Networks.
  - ④ Generative Adversarial Networks.



# Introduction to Deep Neural Networks

Carlos María Alaíz Gudín

---

Overview  
Limitations of MLPs

Introduction

Innovations of Deep Learning

Big Data and Computational Power  
Initialization  
Activation Functions  
Avoiding Over-Fitting  
Specialized Architectures

