

```
!pwd
!ls -la
```

```

/content
total 16
drwxr-xr-x 1 root root 4096 Oct 14 16:31 .
drwxr-xr-x 1 root root 4096 Oct 22 17:56 ..
drwxr-xr-x 1 root root 4096 Oct 14 16:32 .config
drwxr-xr-x 1 root root 4096 Oct 14 16:31 sample_data
```

```
%%writefile labsolReduction0.cu
```

```
// includes, kernels
#include <stdio.h>
#include <assert.h>
```

```
#define NUM_ELEMENTS 512
```

```
// **==-----**
//! @param g_idata  input data in global memory
//               result is expected in index 0 of g_idata
//! @param n        input number of elements to scan from input data
// **==-----**
```

```
__global__ void reduction(float *g_data, int n)
{
    int stride;
    // Define shared memory
    __shared__ float scratch[NUM_ELEMENTS];

    // Load the shared memory
    scratch[threadIdx.x] = g_data[threadIdx.x];
    if(threadIdx.x + blockDim.x < n)
        scratch[threadIdx.x + blockDim.x] = g_data[threadIdx.x + blockDim.x];
    __syncthreads();

    // Do sum reduction from shared memory

    for(stride = 1 ; stride < blockDim.x; stride *= 2)
    {
        __syncthreads();
        if(threadIdx.x % (2*stride) == 0)
            scratch[threadIdx.x] += scratch[threadIdx.x + stride];
    }

    // Store results back to global memory
    if(threadIdx.x == 0)
        g_data[0] = scratch[0];

    return;
}
```

```

////////////////////////////////////
// Program main
////////////////////////////////////
void runTest( int argc, char** argv);
float computeOnDevice(float* h_data, int array_mem_size);
extern "C" void computeGold( float* reference, float* idata, const unsigned int len

int main( int argc, char** argv)
{
    runTest( argc, argv);
    return EXIT_SUCCESS;
}

////////////////////////////////////
//! Run naive scan test
////////////////////////////////////
void runTest( int argc, char** argv)
{
    int num_elements = NUM_ELEMENTS;

    const unsigned int array_mem_size = sizeof( float) * num_elements;

    // allocate host memory to store the input data
    float* h_data = (float*) malloc( array_mem_size);

    // * No arguments: Randomly generate input data and compare against the host's

        // initialize the input data on the host to be integer values
        // between 0 and 1000
        for( unsigned int i = 0; i < num_elements; ++i)
        {
            //h_data[i] = floorf(1000*(rand()/(float)RAND_MAX));
            h_data[i] = i*1.0;

        }

        // compute reference solution
        float reference = 0.0f;
        computeGold(&reference , h_data, num_elements);

        float result = computeOnDevice(h_data, num_elements);

        // We can use an epsilon of 0 since values are integral and in a range
        // that can be exactly represented
        float epsilon = 0.0f;
        unsigned int result_regtest = (abs(result - reference) <= epsilon);
        printf( "Test %s\n", (1 == result_regtest) ? "PASSED" : "FAILED");
        printf( "device: %f  host: %f\n", result, reference);
        // cleanup memory
        free( h_data);
}

////////////////////////////////////

```

```

// Take h_data from host, copies it to device, setup grid and thread
// dimentions, excutes kernel function, and copy result of scan back
// to h_data.
// Note: float* h_data is both the input and the output of this function.
////////////////////////////////////

float computeOnDevice(float* h_data, int num_elements)
{
    float* d_data = NULL;
    float result;

    // Memory allocation on device side
    cudaMalloc((void*)&d_data, num_elements*sizeof(float));

    // Copy from host memory to device memory
    cudaMemcpy(d_data, h_data, num_elements*sizeof(float), cudaMemcpyHostToDevice);

    //int threads = (num_elements/2) + num_elements%2;
    int threads = num_elements;

    // Invoke the kernel
    reduction<<<1,threads>>>(d_data,num_elements);

    // Copy from device memory back to host memory
    cudaMemcpy(&result, d_data, sizeof(float), cudaMemcpyDeviceToHost);
    cudaFree(d_data);
    return result;
}

////////////////////////////////////
void computeGold( float* reference, float* idata, const unsigned int len)
{
    reference[0] = 0;
    double total_sum = 0;
    unsigned int i;
    for( i = 0; i < len; ++i)
    {
        total_sum += idata[i];
    }
    reference[0] = total_sum;
}
////////////////////////////////////

```

➡ Overwriting labsolReduction0.cu

```

!usr/local/cuda/bin/nvcc -I /usr/local/cuda/samples/common/inc -arch=sm_35 -rdc=tr

!./labsolReduction0

```

Test PASSED

!nvprof ./labsolReduction0

==2810== NVPROF is profiling process 2810, command: ./labsolReduction0

Test PASSED

device: 130816.000000 host: 130816.000000

==2810== Profiling application: ./labsolReduction0

==2810== Profiling result:

	Type	Time(%)	Time	Calls	Avg	Min	Max	N
GPU activities:		68.19%	8.0960us	1	8.0960us	8.0960us	8.0960us	r
		16.17%	1.9200us	1	1.9200us	1.9200us	1.9200us	[
		15.63%	1.8560us	1	1.8560us	1.8560us	1.8560us	[
API calls:		99.47%	188.56ms	1	188.56ms	188.56ms	188.56ms	c
		0.30%	572.94us	1	572.94us	572.94us	572.94us	c
		0.13%	246.88us	97	2.5450us	177ns	107.95us	c
		0.04%	77.374us	1	77.374us	77.374us	77.374us	c
		0.02%	37.803us	2	18.901us	17.338us	20.465us	c
		0.02%	29.895us	1	29.895us	29.895us	29.895us	c
		0.01%	27.932us	1	27.932us	27.932us	27.932us	c
		0.00%	4.3930us	1	4.3930us	4.3930us	4.3930us	c
		0.00%	1.8920us	3	630ns	178ns	1.2840us	c
		0.00%	1.2160us	2	608ns	326ns	890ns	c
		0.00%	336ns	1	336ns	336ns	336ns	c

CAS01

%%writefile labsolReduction1.cu

// includes, kernels

#include <stdio.h>

#include <assert.h>

#define NUM_ELEMENTS 512

// **===-----**

//! @param g_idata input data in global memory

// result is expected in index 0 of g_idata

//! @param n input number of elements to scan from input data

// **===-----**

__global__ void reduction(float *g_data, int n)

{

int stride;

// Define shared memory

__shared__ float scratch[NUM_ELEMENTS];

// Load the shared memory

scratch[threadIdx.x] = g_data[threadIdx.x];

if(threadIdx.x + blockDim.x < n)

scratch[threadIdx.x + blockDim.x] = g_data[threadIdx.x + blockDim.x];

__syncthreads();

// Do sum reduction from shared memory

```

for(stride = NUM_ELEMENTS/2 ; stride >= 1; stride >>= 1)
{
    if(threadIdx.x < stride)
        scratch[threadIdx.x] += scratch[threadIdx.x + stride];
    __syncthreads();
}

// Store results back to global memory
if(threadIdx.x == 0)
    g_data[0] = scratch[0];

return;
}

////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////
// Program main
////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////
void runTest( int argc, char** argv);
float computeOnDevice(float* h_data, int array_mem_size);
extern "C" void computeGold( float* reference, float* idata, const unsigned int len

int main( int argc, char** argv)
{
    runTest( argc, argv);
    return EXIT_SUCCESS;
}

////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////
//! Run naive scan test
////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////
void runTest( int argc, char** argv)
{
    int num_elements = NUM_ELEMENTS;

    const unsigned int array_mem_size = sizeof( float) * num_elements;

    // allocate host memory to store the input data
    float* h_data = (float*) malloc( array_mem_size);

    // * No arguments: Randomly generate input data and compare against the host's

        // initialize the input data on the host to be integer values
        // between 0 and 1000
        for( unsigned int i = 0; i < num_elements; ++i)
        {
            //h_data[i] = floorf(1000*(rand()/(float)RAND_MAX));
            h_data[i] = i*1.0;
        }

        // compute reference solution
        float reference = 0.0f;
        computeGold(&reference , h_data, num_elements);

```

```

float result = computeOnDevice(h_data, num_elements);

// We can use an epsilon of 0 since values are integral and in a range
// that can be exactly represented
float epsilon = 0.0f;
unsigned int result_regtest = (abs(result - reference) <= epsilon);
printf( "Test %s\n", (1 == result_regtest) ? "PASSED" : "FAILED");
printf( "device: %f  host: %f\n", result, reference);
// cleanup memory
free( h_data);
}

////////////////////////////////////
// Take h_data from host, copies it to device, setup grid and thread
// dimentions, excutes kernel function, and copy result of scan back
// to h_data.
// Note: float* h_data is both the input and the output of this function.
////////////////////////////////////

float computeOnDevice(float* h_data, int num_elements)
{
    float* d_data = NULL;
    float result;

    // Memory allocation on device side
    cudaMalloc((void*)&d_data, num_elements*sizeof(float));

    // Copy from host memory to device memory
    cudaMemcpy(d_data, h_data, num_elements*sizeof(float), cudaMemcpyHostToDevice);

    int threads = (num_elements/2) + num_elements%2;

    // Invoke the kernel
    reduction<<<1,threads>>>(d_data,num_elements);

    // Copy from device memory back to host memory
    cudaMemcpy(&result, d_data, sizeof(float), cudaMemcpyDeviceToHost);
    cudaFree(d_data);
    return result;
}

////////////////////////////////////
void computeGold( float* reference, float* idata, const unsigned int len)
{
    reference[0] = 0;
    double total_sum = 0;
    unsigned int i;
    for( i = 0; i < len; ++i)
    {
        total_sum += idata[i];
    }
    reference[0] = total_sum;
}

////////////////////////////////////

```

//

Writing labsolReduction1.cu

!/usr/local/cuda/bin/nvcc -I /usr/local/cuda/samples/common/inc -arch=sm_35 -rdc=tr

!ls

lab1Reduction lab1solReduction labsolReduction1 sample_data
lab1Reduction.cu lab1solReduction.cu labsolReduction1.cu

!./labsolReduction1

Test PASSED
device: 523776.000000 host: 523776.000000

!nvprof ./labsolReduction1

==1123== NVPROF is profiling process 1123, command: ./labsolReduction1
Test PASSED
device: 523776.000000 host: 523776.000000
==1123== Profiling application: ./labsolReduction1
==1123== Profiling result:

	Type	Time(%)	Time	Calls	Avg	Min	Max	N
GPU activities:		56.94%	5.1200us	1	5.1200us	5.1200us	5.1200us	r
		22.78%	2.0480us	1	2.0480us	2.0480us	2.0480us	[
		20.28%	1.8240us	1	1.8240us	1.8240us	1.8240us	[
API calls:		99.63%	196.11ms	1	196.11ms	196.11ms	196.11ms	c
		0.20%	398.73us	1	398.73us	398.73us	398.73us	c
		0.08%	153.97us	97	1.5870us	143ns	59.096us	c
		0.04%	77.037us	1	77.037us	77.037us	77.037us	c
		0.02%	35.271us	2	17.635us	14.457us	20.814us	c
		0.01%	27.427us	1	27.427us	27.427us	27.427us	c
		0.01%	24.317us	1	24.317us	24.317us	24.317us	c
		0.00%	3.3170us	1	3.3170us	3.3170us	3.3170us	c
		0.00%	2.0920us	3	697ns	162ns	1.3160us	c
		0.00%	1.4470us	2	723ns	304ns	1.1430us	c
		0.00%	269ns	1	269ns	269ns	269ns	c

CAS02

%%writefile labsolReduction2.cu

// includes, kernels
#include <stdio.h>
#include <assert.h>

#define NUM_ELEMENTS 512

```

// **==-----**
//! @param g_idata  input data in global memory
//          result is expected in index 0 of g_idata
//! @param n        input number of elements to scan from input data
// **==-----**

__global__ void reduction(float *g_data, int n)
{
    int stride;
    // Define shared memory
    __shared__ float scratch[NUM_ELEMENTS];

    // Load the shared memory
    scratch[threadIdx.x] = g_data[threadIdx.x];
    if(threadIdx.x + blockDim.x < n)
        scratch[threadIdx.x + blockDim.x] = g_data[threadIdx.x + blockDim.x];
    __syncthreads();

    // Do sum reduction from shared memory
    //Reduction scheme 2
    for (stride = NUM_ELEMENTS; stride > 1; stride >>= 1) {
        __syncthreads();
        if (threadIdx.x < (stride>>1)) {
            scratch[threadIdx.x] += scratch[stride - threadIdx.x - 1];
        }
    }
    __syncthreads();

    // Store results back to global memory
    if(threadIdx.x == 0)
        g_data[0] = scratch[0];

    return;
}

////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////
// Program main
////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////
void runTest( int argc, char** argv);
float computeOnDevice(float* h_data, int array_mem_size);
extern "C" void computeGold( float* reference, float* idata, const unsigned int len

int main( int argc, char** argv)
{
    runTest( argc, argv);
    return EXIT_SUCCESS;
}

////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////
//! Run naive scan test
////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////
void runTest( int argc, char** argv)
{
    int num_elements = NUM_ELEMENTS;

```



```

const unsigned int array_mem_size = sizeof( float) * num_elements;

// allocate host memory to store the input data
float* h_data = (float*) malloc( array_mem_size);

// * No arguments: Randomly generate input data and compare against the host's

    // initialize the input data on the host to be integer values
    // between 0 and 1000
    for( unsigned int i = 0; i < num_elements; ++i)
    {
        //h_data[i] = floorf(1000*(rand()/(float)RAND_MAX));
        h_data[i] = i*1.0;
    }

    // compute reference solution
    float reference = 0.0f;
    computeGold(&reference , h_data, num_elements);

    float result = computeOnDevice(h_data, num_elements);

    // We can use an epsilon of 0 since values are integral and in a range
    // that can be exactly represented
    float epsilon = 0.0f;
    unsigned int result_regtest = (abs(result - reference) <= epsilon);
    printf( "Test %s\n", (1 == result_regtest) ? "PASSED" : "FAILED");
    printf( "device: %f  host: %f\n", result, reference);
    // cleanup memory
    free( h_data);
}

////////////////////////////////////
// Take h_data from host, copies it to device, setup grid and thread
// dimentions, excutes kernel function, and copy result of scan back
// to h_data.
// Note: float* h_data is both the input and the output of this function.
////////////////////////////////////

float computeOnDevice(float* h_data, int num_elements)
{
    float* d_data = NULL;
    float result;

    // Memory allocation on device side
    cudaMalloc((void*)&d_data, num_elements*sizeof(float));

    // Copy from host memory to device memory
    cudaMemcpy(d_data, h_data, num_elements*sizeof(float), cudaMemcpyHostToDevice);

    int threads = (num_elements/2) + num_elements%2;

    // Invoke the kernel
    reduction<<<1,threads>>>(d_data,num_elements);

```

```

// Copy from device memory back to host memory
cudaMemcpy(&result, d_data, sizeof(float), cudaMemcpyDeviceToHost);
cudaFree(d_data);
return result;
}

////////////////////////////////////
void computeGold( float* reference, float* idata, const unsigned int len)
{
    reference[0] = 0;
    double total_sum = 0;
    unsigned int i;
    for( i = 0; i < len; ++i)
    {
        total_sum += idata[i];
    }
    reference[0] = total_sum;
}
////////////////////////////////////

```

Overwriting labsolReduction2.cu

```
!usr/local/cuda/bin/nvcc -I /usr/local/cuda/samples/common/inc -arch=sm_35 -rdc=tr
```

```
!ls -la
```

```

total 3280
drwxr-xr-x 1 root root 4096 Oct 22 19:54 .
drwxr-xr-x 1 root root 4096 Oct 22 18:46 ..
drwxr-xr-x 1 root root 4096 Oct 14 16:32 .config
-rwxr-xr-x 1 root root 659784 Oct 22 18:57 lab1Reduction
-rw-r--r-- 1 root root 6837 Oct 22 18:54 lab1Reduction.cu
-rwxr-xr-x 1 root root 655504 Oct 22 19:43 lab1solReduction
-rw-r--r-- 1 root root 5121 Oct 22 19:30 lab1solReduction.cu
-rwxr-xr-x 1 root root 655504 Oct 22 19:37 labsolReduction1
-rw-r--r-- 1 root root 4338 Oct 22 19:35 labsolReduction1.cu
-rwxr-xr-x 1 root root 655504 Oct 22 19:54 labsolReduction2
-rw-r--r-- 1 root root 4371 Oct 22 19:45 labsolReduction2.cu
-rwxr-xr-x 1 root root 655504 Oct 22 19:50 labsolReduction3
-rw-r--r-- 1 root root 4408 Oct 22 19:49 labsolReduction3.cu
drwxr-xr-x 1 root root 4096 Oct 14 16:31 sample_data

```

```
!./labsolReduction2
```

```

Test FAILED
device: 81760.000000 host: 130816.000000

```

```
!nvprof ./labsolReduction2
```

==1931== NVPROF is profiling process 1931, command: ./labsolReduction2

Test PASSED

device: 130816.000000 host: 130816.000000

==1931== Profiling application: ./labsolReduction2

==1931== Profiling result:

	Type	Time(%)	Time	Calls	Avg	Min	Max	N
GPU activities:		53.54%	4.3520us	1	4.3520us	4.3520us	4.3520us	r
		23.62%	1.9200us	1	1.9200us	1.9200us	1.9200us	[
		22.83%	1.8560us	1	1.8560us	1.8560us	1.8560us	[
API calls:		99.67%	228.07ms	1	228.07ms	228.07ms	228.07ms	c
		0.17%	380.32us	1	380.32us	380.32us	380.32us	c
		0.06%	145.35us	97	1.4980us	151ns	60.700us	c
		0.05%	114.10us	1	114.10us	114.10us	114.10us	c
		0.02%	52.420us	2	26.210us	19.272us	33.148us	c
		0.01%	27.815us	1	27.815us	27.815us	27.815us	c
		0.01%	26.198us	1	26.198us	26.198us	26.198us	c
		0.00%	3.5080us	1	3.5080us	3.5080us	3.5080us	c

CASO3

```
%%writefile labsolReduction3.cu
```

```
// includes, kernels
```

```
#include <stdio.h>
```

```
#include <assert.h>
```

```
#define NUM_ELEMENTS 512
```

```
// **====-----*****
//! @param g_idata  input data in global memory
//          result is expected in index 0 of g_idata
//! @param n        input number of elements to scan from input data
// **====-----*****
```

```
__global__ void reduction(float *g_data, int n)
```

```
{
    int stride;
    // Define shared memory
    __shared__ float scratch[NUM_ELEMENTS];

    // Load the shared memory
    scratch[threadIdx.x] = g_data[threadIdx.x];
    if(threadIdx.x + blockDim.x < n)
        scratch[threadIdx.x + blockDim.x] = g_data[threadIdx.x + blockDim.x];
    __syncthreads();
}
```

```
// Do sum reduction from shared memory
```

```
//////// Reduction scheme 3 //////////
```

```
int stride2;
for (stride = 1, stride2=1; stride <= 9; stride++, stride2<=<=1)
{
    int t = threadIdx.x << stride;
    if (t + stride2 < n)
        scratch[t] = scratch[t] + scratch[t + stride2];
}
```

```

    __syncthreads();
}

// Store results back to global memory
if(threadIdx.x == 0)
    g_data[0] = scratch[0];

return;
}
/////////////////////////////////////////////////////////////////
// Program main
/////////////////////////////////////////////////////////////////
void runTest( int argc, char** argv);
float computeOnDevice(float* h_data, int array_mem_size);
extern "C" void computeGold( float* reference, float* idata, const unsigned int len

int main( int argc, char** argv)
{
    runTest( argc, argv);
    return EXIT_SUCCESS;
}

/////////////////////////////////////////////////////////////////
//! Run naive scan test
/////////////////////////////////////////////////////////////////
void runTest( int argc, char** argv)
{
    int num_elements = NUM_ELEMENTS;

    const unsigned int array_mem_size = sizeof( float) * num_elements;

    // allocate host memory to store the input data
    float* h_data = (float*) malloc( array_mem_size);

    // * No arguments: Randomly generate input data and compare against the host's

        // initialize the input data on the host to be integer values
        // between 0 and 1000
        for( unsigned int i = 0; i < num_elements; ++i)
        {
            //h_data[i] = floorf(1000*(rand()/(float)RAND_MAX));
            h_data[i] = i*1.0;
        }

        // compute reference solution
        float reference = 0.0f;
        computeGold(&reference , h_data, num_elements);

        float result = computeOnDevice(h_data, num_elements);

        // We can use an epsilon of 0 since values are integral and in a range
        // that can be exactly represented
        float epsilon = 0.0f;
        unsigned int result_regtest = (abs(result - reference) <= epsilon);

```

```

    printf( "Test %s\n", (1 == result_regtest) ? "PASSED" : "FAILED");
    printf( "device: %f  host: %f\n", result, reference);
    // cleanup memory
    free( h_data);
}

////////////////////////////////////
// Take h_data from host, copies it to device, setup grid and thread
// dimentions, excutes kernel function, and copy result of scan back
// to h_data.
// Note: float* h_data is both the input and the output of this function.
////////////////////////////////////

float computeOnDevice(float* h_data, int num_elements)
{
    float* d_data = NULL;
    float result;

    // Memory allocation on device side
    cudaMalloc((void*)&d_data, num_elements*sizeof(float));

    // Copy from host memory to device memory
    cudaMemcpy(d_data, h_data, num_elements*sizeof(float), cudaMemcpyHostToDevice);

    int threads = (num_elements/2) + num_elements%2;

    // Invoke the kernel
    reduction<<<1,threads>>>(d_data,num_elements);

    // Copy from device memory back to host memory
    cudaMemcpy(&result, d_data, sizeof(float), cudaMemcpyDeviceToHost);
    cudaFree(d_data);
    return result;
}

////////////////////////////////////
void computeGold( float* reference, float* idata, const unsigned int len)
{
    reference[0] = 0;
    double total_sum = 0;
    unsigned int i;
    for( i = 0; i < len; ++i)
    {
        total_sum += idata[i];
    }
    reference[0] = total_sum;
}
////////////////////////////////////

```

Overwriting labsolReduction3.cu

```
! /usr/local/cuda/bin/nvcc -I /usr/local/cuda/samples/common/inc -arch=sm_35 -rdc=tr
```

```
!ls -la
```

```
!./labsolReduction3
```

```
Test PASSED
device: 130816.000000 host: 130816.000000
```

```
!nvprof ./labsolReduction3
```

```
==2116== NVPROF is profiling process 2116, command: ./labsolReduction3
```

```
Test PASSED
```

```
device: 130816.000000 host: 130816.000000
```

```
==2116== Profiling application: ./labsolReduction3
```

```
==2116== Profiling result:
```

	Type	Time(%)	Time	Calls	Avg	Min	Max	N
GPU activities:		54.41%	4.5440us	1	4.5440us	4.5440us	4.5440us	r
		23.37%	1.9520us	1	1.9520us	1.9520us	1.9520us	[
		22.22%	1.8560us	1	1.8560us	1.8560us	1.8560us	[
API calls:		99.62%	191.23ms	1	191.23ms	191.23ms	191.23ms	c
		0.20%	378.14us	1	378.14us	378.14us	378.14us	c
		0.09%	166.79us	97	1.7190us	157ns	67.113us	c
		0.04%	78.328us	1	78.328us	78.328us	78.328us	c
		0.02%	35.975us	2	17.987us	15.838us	20.137us	c
		0.02%	29.526us	1	29.526us	29.526us	29.526us	c
		0.01%	27.831us	1	27.831us	27.831us	27.831us	c
		0.00%	3.6840us	1	3.6840us	3.6840us	3.6840us	c
		0.00%	1.7880us	3	596ns	155ns	1.2790us	c
		0.00%	1.0180us	2	509ns	326ns	692ns	c
		0.00%	327ns	1	327ns	327ns	327ns	c

```
!nvprof ./labsolReduction2
```

```
==2092== NVPROF is profiling process 2092, command: ./labsolReduction2
```

```
Test PASSED
```

```
device: 130816.000000 host: 130816.000000
```

```
==2092== Profiling application: ./labsolReduction2
```

```
==2092== Profiling result:
```

	Type	Time(%)	Time	Calls	Avg	Min	Max	N
GPU activities:		53.75%	4.3520us	1	4.3520us	4.3520us	4.3520us	r
		23.32%	1.8880us	1	1.8880us	1.8880us	1.8880us	[
		22.92%	1.8560us	1	1.8560us	1.8560us	1.8560us	[
API calls:		99.64%	194.61ms	1	194.61ms	194.61ms	194.61ms	c
		0.19%	364.12us	1	364.12us	364.12us	364.12us	c
		0.08%	161.47us	97	1.6640us	150ns	71.381us	c
		0.04%	77.930us	1	77.930us	77.930us	77.930us	c
		0.02%	38.405us	2	19.202us	18.295us	20.110us	c
		0.02%	29.673us	1	29.673us	29.673us	29.673us	c
		0.01%	26.666us	1	26.666us	26.666us	26.666us	c
		0.00%	3.3290us	1	3.3290us	3.3290us	3.3290us	c
		0.00%	2.3010us	3	767ns	165ns	1.6380us	c
		0.00%	1.2510us	2	625ns	342ns	909ns	c
		0.00%	330ns	1	330ns	330ns	330ns	c

```
!nvprof ./labsolReduction1
```

==2103== NVPROF is profiling process 2103, command: ./labsolReduction1

Test PASSED

device: 523776.000000 host: 523776.000000

==2103== Profiling application: ./labsolReduction1

==2103== Profiling result:

Type	Time(%)	Time	Calls	Avg	Min	Max	N
GPU activities:	57.04%	5.1840us	1	5.1840us	5.1840us	5.1840us	r
	22.89%	2.0800us	1	2.0800us	2.0800us	2.0800us	[
	20.07%	1.8240us	1	1.8240us	1.8240us	1.8240us	[
API calls:	99.63%	204.88ms	1	204.88ms	204.88ms	204.88ms	c
	0.19%	398.08us	1	398.08us	398.08us	398.08us	c
	0.07%	149.94us	97	1.5450us	155ns	63.211us	c
	0.05%	95.307us	1	95.307us	95.307us	95.307us	c
	0.02%	48.822us	2	24.411us	16.850us	31.972us	c
	0.01%	27.708us	1	27.708us	27.708us	27.708us	c
	0.01%	26.020us	1	26.020us	26.020us	26.020us	c
	0.00%	3.4330us	1	3.4330us	3.4330us	3.4330us	c
	0.00%	1.8310us	3	610ns	176ns	1.2030us	c
	0.00%	1.3370us	2	668ns	259ns	1.0780us	c
	0.00%	318ns	1	318ns	318ns	318ns	c