

# Procesamiento de Señal y Transformadas

## Convolutional Neural Networks

Prof. Rubén Vera Rodríguez

ruben.vera@uam.es

BiDA Lab, EPS

<http://atvs.ii.uam.es/atvs/>

Based on material by Fei-Fei Li & Justin Johnson & Serena Yeung from Stanford University

# Bit of history...

Frank Rosenblatt, ~1957: Perceptron


Widrow and Hoff, ~1960: Adaline/Madaline: Multilayer perceptron networks

Rumelhart et al., 1986: First time back-propagation became popular

# Bit of history...

## Convolutional Networks: 1998

LeNet: a layered model composed of convolution and subsampling operations followed by a holistic representation and ultimately a classifier for handwritten digits. [LeCun, Bottou, Bengio, Haffner 1998]



# Bit of history...

Hinton and Salakhutdinov 2006, Reinvigorated research in Deep Learning

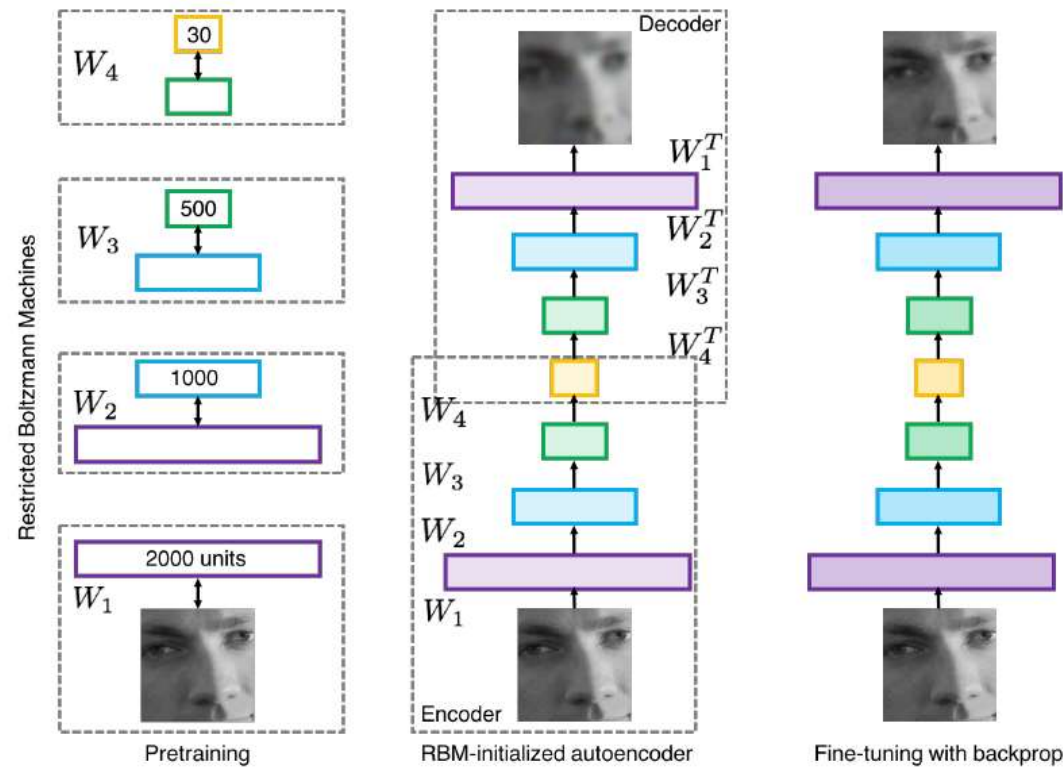



Illustration of Hinton and Salakhutdinov 2006 by Lane McIntosh, copyright CS231n 2017

# Bit of history...

## Convolutional Nets: 2012

AlexNet: a layered model composed of convolution, subsampling, and further operations followed by a holistic representation and a final classification on ILSVRC12.

+ data + gpu + non-saturating nonlinearity + regularization



ImageNet Classification with Deep Convolutional Neural Networks [Krizhevsky, Sutskever, Hinton, 2012]

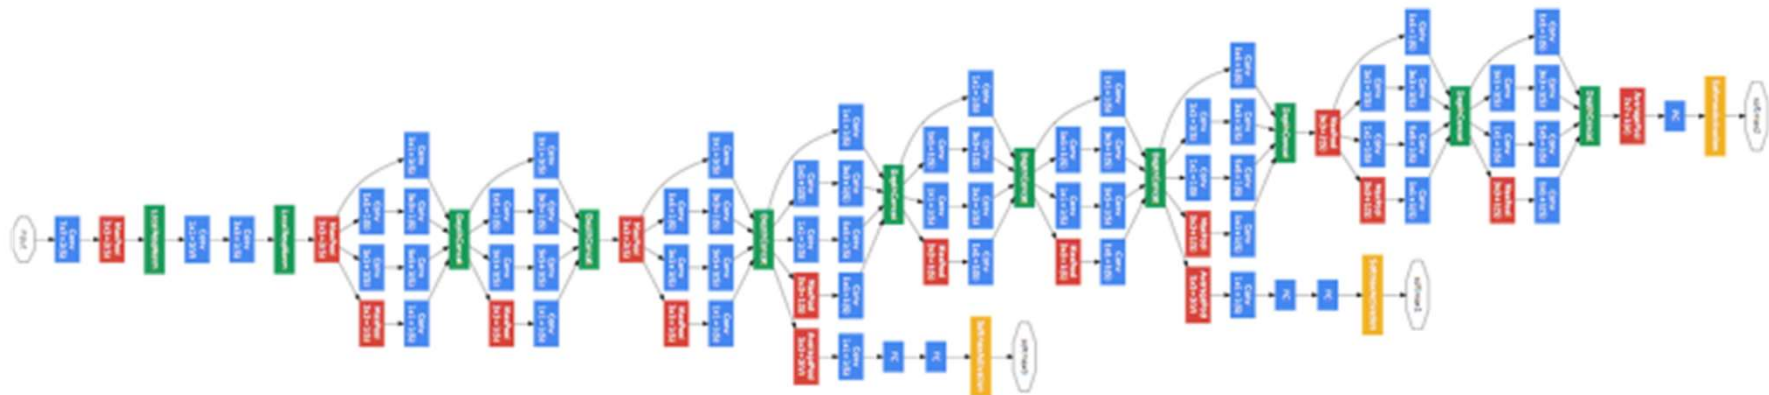
# CNN from LeNet to GoogLeNet

## Convolutional Nets: 2014


**ILSVRC**14 Winners: ~6.6% Top-5 error

- **GoogLeNet**: composition of multi-scale dimension-reduced modules (pictured)
- **VGG**: 16 layers of 3x3 convolution interleaved with max pooling + 3 fully-connected layers

+ data + gpu + non-saturating nonlinearity + regularization



# Evolution of Neural Networks



1. Perceptron (Rosenblatt, 1958, 1962)
2. Adaptive linear element (Widrow and Hoff, 1960)
3. Neocognitron (Fukushima, 1980)
4. Early back-propagation network (Rumelhart *et al.*, 1986b)
5. Recurrent neural network for speech recognition (Robinson and Fallside, 1991)
6. Multilayer perceptron for speech recognition (Bengio *et al.*, 1991)
7. Mean field sigmoid belief network (Saul *et al.*, 1996)
8. LeNet-5 (LeCun *et al.*, 1998b)
9. Echo state network (Jaeger and Haas, 2004)
10. Deep belief network (Hinton *et al.*, 2006)
11. GPU-accelerated convolutional network (Chellapilla *et al.*, 2006)
12. Deep Boltzmann machine (Salakhutdinov and Hinton, 2009a)
13. GPU-accelerated deep belief network (Raina *et al.*, 2009)
14. Unsupervised convolutional network (Jarrett *et al.*, 2009)
15. GPU-accelerated multilayer perceptron (Ciresan *et al.*, 2010)
16. OMP-1 network (Coates and Ng, 2011)
17. Distributed autoencoder (Le *et al.*, 2012)
18. Multi-GPU convolutional network (Krizhevsky *et al.*, 2012)
19. COTS HPC unsupervised convolutional network (Coates *et al.*, 2013)
20. GoogLeNet (Szegedy *et al.*, 2014a)



# ConvNets Applications

Classification



Retrieval



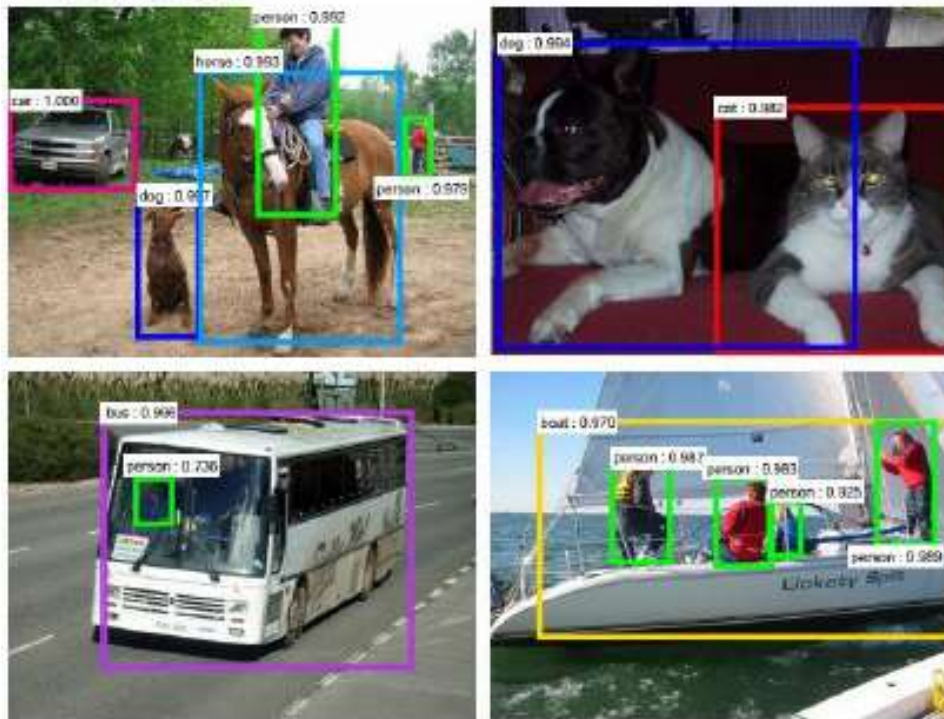
Figures copyright Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, 2012. Reproduced with permission.

ImageNet Classification with Deep Convolutional Neural Networks [Krizhevsky, Sutskever, Hinton, 2012]



# ConvNets Applications


## Detection



Figures copyright Shaoqing Ren, Kaiming He, Ross Girshick, Jian Sun, 2015. Reproduced with permission.

[Faster R-CNN: Ren, He, Girshick, Sun 2015]

## Segmentation



Figures copyright Clement Farabet, 2012. Reproduced with permission.

[Farabet et al., 2012]

# ConvNets Applications



Photo by Lane McIntosh. Copyright CS231n 2017.

self-driving cars



[This image](#) by GBPublic\_PR is licensed under [CC-BY 2.0](#)

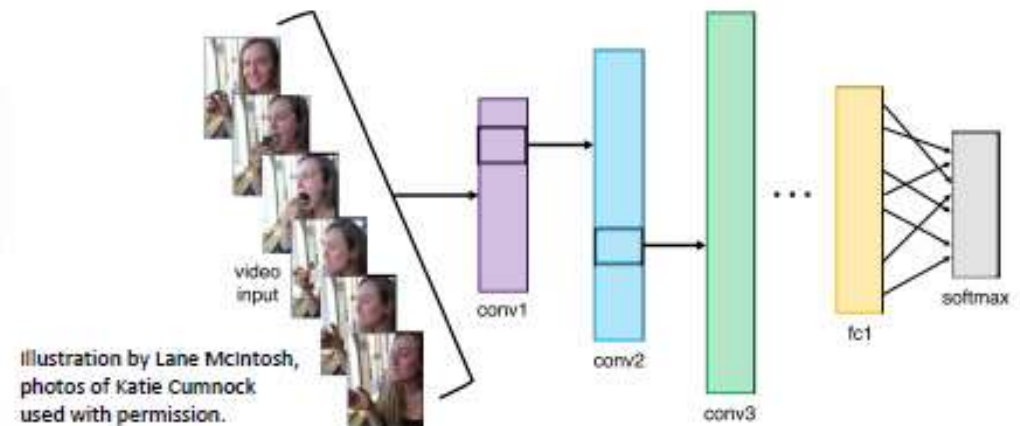
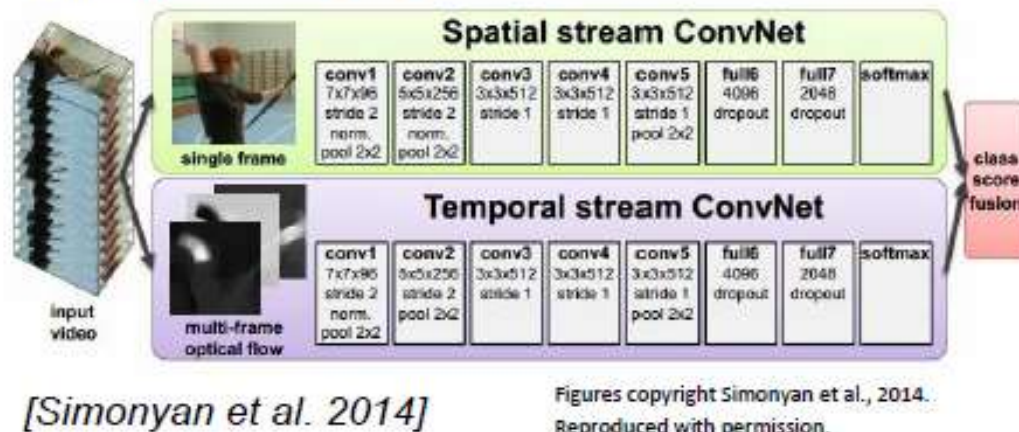

## NVIDIA Tesla line

(these are the GPUs on rye01.stanford.edu)

Note that for embedded systems a typical setup would involve NVIDIA Tegras, with integrated GPU and ARM-based CPU cores.



# ConvNets Applications




# ConvNets Applications

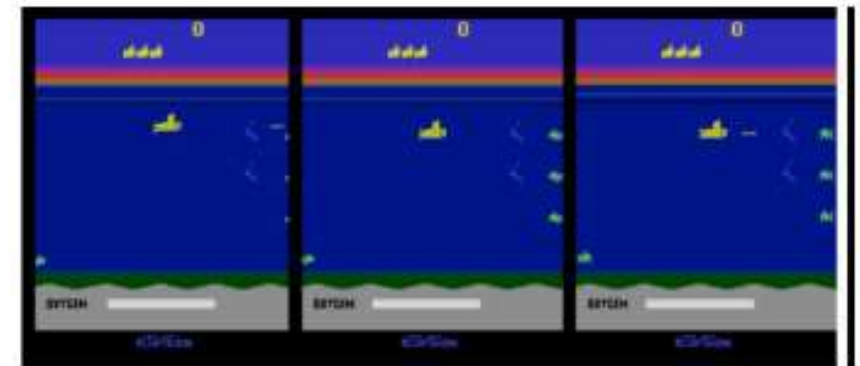


Images are examples of pose estimation, not actually from Toshev & Szegedy 2014. Copyright Lane McIntosh.

[Toshev, Szegedy 2014]



[Guo et al. 2014]



Figures copyright Xiaoxiao Guo, Satinder Singh, Honglak Lee, Richard Lewis, and Xiaoshi Wang, 2014. Reproduced with permission.

# How can we extract features from unstructured data?

Image Domain



Translator



Feature Domain

$$\mathbf{f} = \begin{bmatrix} f_1 \\ f_2 \\ \vdots \\ f_{128} \end{bmatrix}$$




# How can we extract features from unstructured data?

Image Domain




Deep Learning algorithm



Feature Domain


$$\mathbf{f} = \begin{bmatrix} f_1 \\ f_2 \\ \vdots \\ f_{128} \end{bmatrix}$$

# Fully Connected Layer




32x32x3 image -> stretch to 3072 x 1


# Fully Connected Layer




32x32x3 image -> stretch to 3072 x 1



# Fully Connected Layer



32x32x3 image -> stretch to 3072 x 1




**1 number:**

the result of taking a dot product  
between a row of  $W$  and the input  
(a 3072-dimensional dot product)

# Conv Layer


32x32x3 image -> preserve spatial structure





# Conv Layer

32x32x3 image -> preserve spatial structure




5x5x3 filter



**Convolve** the filter with the image  
i.e. “slide over the image spatially,  
computing dot products”

# Conv Layer

32x32x3 image -> preserve spatial structure




Filters always extend the full depth of the input volume

5x5x3 filter



**Convolve** the filter with the image  
i.e. “slide over the image spatially,  
computing dot products”


# Conv Layer



**1 number:**

the result of taking a dot product between the filter and a small  $5 \times 5 \times 3$  chunk of the image (i.e.  $5 \times 5 \times 3 = 75$ -dimensional dot product + bias)


# Conv Layer




**1 number:**

the result of taking a dot product between the filter and a small 5x5x3 chunk of the image (i.e.  $5 \times 5 \times 3 = 75$ -dimensional dot product + bias)

The filter is just a neuron with local connectivity...




# Conv Layer






# Conv Layer

Consider a second filter



# Conv Layer


For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:




We stack these up to get a “new image” of size 28x28x6!

# ConvNet

**Preview:** ConvNet is a sequence of Convolution Layers, interspersed with activation functions




# Spatial Dimensions



# Spatial Dimensions

A closer look at spatial dimensions:




7x7 input (spatially)  
assume 3x3 filter



# Spatial Dimensions

A closer look at spatial dimensions:




7x7 input (spatially)  
assume 3x3 filter

=> **5x5 output**

# Spatial Dimensions


A closer look at spatial dimensions:



7x7 input (spatially)  
assume 3x3 filter  
applied **with stride 2**

# Spatial Dimensions


A closer look at spatial dimensions:



7x7 input (spatially)  
assume 3x3 filter  
applied **with stride 2**  
**=> 3x3 output!**

# Spatial Dimensions


A closer look at spatial dimensions:



7x7 input (spatially)  
assume 3x3 filter  
applied **with stride 3?**

# Spatial Dimensions


A closer look at spatial dimensions:



7x7 input (spatially)  
assume 3x3 filter  
applied **with stride 3?**

**doesn't fit!**  
cannot apply 3x3 filter on  
7x7 input with stride 3.

# Spatial Dimensions



Output size:  
 $(N - F) / \text{stride} + 1$

e.g.  $N = 7, F = 3$ :

stride 1  $\Rightarrow (7 - 3) / 1 + 1 = 5$

stride 2  $\Rightarrow (7 - 3) / 2 + 1 = 3$

stride 3  $\Rightarrow (7 - 3) / 3 + 1 = 2.33 : \backslash$

# Spatial Dimensions

In practice: Common to zero pad the border

0	0	0	0	0	0			
0								
0								
0								
0								

e.g. input 7x7

**3x3** filter, applied with **stride 1**

**pad with 1 pixel** border => what is the output?

(recall:)

$$(N - F) / \text{stride} + 1$$



# Spatial Dimensions

In practice: Common to zero pad the border

0	0	0	0	0	0			
0								
0								
0								
0								

e.g. input 7x7

**3x3** filter, applied with **stride 1**

**pad with 1 pixel** border => what is the output?

**7x7 output!**

(recall:)

$$(N + 2P - F) / \text{stride} + 1$$

# Spatial Dimensions

In practice: Common to zero pad the border

0	0	0	0	0	0			
0								
0								
0								
0								

e.g. input 7x7

**3x3** filter, applied with **stride 1**

**pad with 1 pixel** border => what is the output?

**7x7 output!**

in general, common to see CONV layers with stride 1, filters of size  $F \times F$ , and zero-padding with  $(F-1)/2$ . (will preserve size spatially)

e.g.  $F = 3 \Rightarrow$  zero pad with 1


$F = 5 \Rightarrow$  zero pad with 2

$F = 7 \Rightarrow$  zero pad with 3

# Spatial Dimensions

**Remember back to...**

E.g. 32x32 input convolved repeatedly with 5x5 filters shrinks volumes spatially! (32 -> 28 -> 24 ...). Shrinking too fast is not good, doesn't work well.




# Spatial Dimensions

Examples time:

Input volume: **32x32x3**

10 5x5 filters with stride 1, pad 2

Output volume size: ?




# Spatial Dimensions

Examples time:

Input volume: **32x32x3**

10 5x5 filters with stride 1, pad 2



Output volume size: ?

$(32 + 2 * 2 - 5) / 1 + 1 = 32$  spatially, so

**32x32x10**




# Spatial Dimensions

Examples time:

Input volume: **32x32x3**

10 5x5 filters with stride 1, pad 2




Number of parameters in this layer?

# Spatial Dimensions

Examples time:

Input volume: **32x32x3**

10 5x5 filters with stride 1, pad 2




Number of parameters in this layer?

each filter has  $5*5*3 + 1 = 76$  params (+1 for bias)

$\Rightarrow 76*10 = 760$

# Spatial Dimensions



Input volume=5x5x3  
Filter 3x3x3  
Stride =2

Output

o[:, :, 0]		
3	-2	-5
-1	-2	-5
-4	-8	-4

# Spatial Dimensions

## Convolution layer: summary

Let's assume input is  $W_1 \times H_1 \times C$

Conv layer needs 4 hyperparameters:

- Number of filters **K**
- The filter size **F**
- The stride **S**
- The zero padding **P**

This will produce an output of  $W_2 \times H_2 \times K$   
where:

- $W_2 = (W_1 - F + 2P)/S + 1$
- $H_2 = (H_1 - F + 2P)/S + 1$

Number of parameters:  $F^2CK$  and  $K$  biases



# Spatial Dimensions

## Convolution layer: summary

Let's assume input is  $W_1 \times H_1 \times C$

Conv layer needs 4 hyperparameters:

- Number of filters **K**
- The filter size **F**
- The stride **S**
- The zero padding **P**

This will produce an output of  $W_2 \times H_2 \times K$   
where:

- $W_2 = (W_1 - F + 2P)/S + 1$
- $H_2 = (H_1 - F + 2P)/S + 1$

Number of parameters:  $F^2CK$  and  $K$  biases

Common settings:


$K =$  (powers of 2, e.g. 32, 64, 128, 512)

- $F = 3, S = 1, P = 1$
- $F = 5, S = 1, P = 2$
- $F = 5, S = 2, P = ?$  (whatever fits)
- $F = 1, S = 1, P = 0$



# Spatial Dimensions


(btw, 1x1 convolution layers make perfect sense)



# Conv Layer

An activation map is a 28x28 sheet of neuron outputs:

1. Each is connected to a small region in the input
2. All of them share parameters




“5x5 filter” -> “5x5 receptive field for each neuron”

# Conv Layer

E.g. with 5 filters, CONV layer consists of neurons arranged in a 3D grid (28x28x5)


There will be 5 different neurons all looking at the same region in the input volume



# Fully Connected Layer

32x32x3 image -> stretch to 3072 x 1


Each neuron looks at the full input volume



**1 number:**  
the result of taking a dot product  
between a row of  $W$  and the input  
(a 3072-dimensional dot product)


# ConvNet

**Preview:** ConvNet is a sequence of Convolution Layers, interspersed with activation functions







# ConvNet




# Activation Functions

$f$  is the activation function:


These are non-linear functions applied at the output values



Sigmoid



Tanh




ReLu

Very popular, simple implementation


# Pooling Layer

Makes the representations smaller and more manageable

Operates over each activation map independently:




# Pooling Layer: Max Pooling




# Fully Connected Layer

Contains neurons that connect to the entire input volume, as in ordinary Neural Networks



# Convolutional Neural Networks: AlexNet

Slides from Jason Yosinski




Krizhevsky, Sutskever, Hinton – NIPS 2012



# Convolutional Neural Networks: AlexNet


Slides from Jason Yosinski




Krizhevsky, Sutskever, Hinton – NIPS 2012

# Convolutional Neural Networks: AlexNet

Slides from Jason Yosinski



Layer 1  
Filter  
(Edges and  
color blobs)




# Convolutional Neural Networks: AlexNet

Slides from Jason Yosinski



# Convolutional Neural Networks: AlexNet

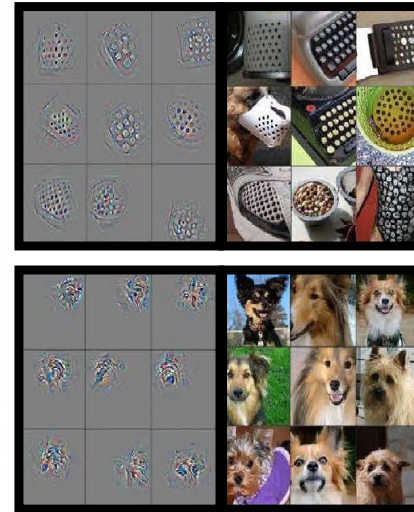
Slides from Jason Yosinski



Layer 1  
Filter  
(Edges and  
color blobs)




Layer 2



Layer 5



Windsor tie: 0.998959



Windsor tie: 0.992462


Last  
Layer

Zeiler et al.  
arXiv 2013

Nguyen et al.  
arXiv 2014


# Convolutional Neural Networks: AlexNet

Slides from Jason Yosinski



# Convolutional Neural Networks: AlexNet


Slides from Jason Yosinski






# Convolutional Neural Networks: AlexNet

Slides from Jason Yosinski




# Convolutional Neural Networks: AlexNet

Slides from Jason Yosinski




# Convolutional Neural Networks: AlexNet

Slides from Jason Yosinski



# Transfer learning



General purpose pre-trained network




Inception from Google

Weights adapted to fashion images




Fashion-net

# Transfer learning: embeddings



**Tuned network  
for fashion  
applications**

# Transfer learning: embeddings






# Transfer learning

## Transfer Learning with CNNs

1. Train on Imagenet




2. Small Dataset (C classes)




Donahue et al, "DeCAF: A Deep Convolutional Activation Feature for Generic Visual Recognition", ICML 2014  
Razavian et al, "CNN Features Off-the-Shelf: An Astounding Baseline for Recognition", CVPR Workshops 2014

Finetuned from AlexNet



Donahue et al, "DeCAF: A Deep Convolutional Activation Feature for Generic Visual Recognition", ICML 2014

# Transfer learning



More specific

More generic


	very similar dataset	very different dataset
very little data	Finetune linear classifier on top layer	You're in trouble... Try data augmentation / collect more data
quite a lot of data	Finetune a few layers	Finetune a larger number of layers

# Data Augmentation

a. No augmentation (= 1 image)




b. Flip augmentation (= 2 images)




c. Crop+Flip augmentation (= 10 images)




# Dropout



(a) Standard Neural Net



(b) After applying dropout.




Dropout: A Simple Way to Prevent Neural Networks from Overfitting  
(N Srivastava et al. 2015)

# How can we extract features from unstructured data?

Image Domain



Deep Learning algorithm



Feature Domain

$$\mathbf{f} = \begin{bmatrix} f_1 \\ f_2 \\ \vdots \\ f_{128} \end{bmatrix}$$



# We can compare images in the feature domain



$$\mathbf{f}^1 = \begin{bmatrix} f_1^1 \\ f_2^1 \\ \vdots \\ f_{128}^1 \end{bmatrix}$$




$$\mathbf{f}^1 = \begin{bmatrix} f_1^1 \\ f_2^1 \\ \vdots \\ f_{128}^1 \end{bmatrix}$$

$$d(\mathbf{f}^1, \mathbf{f}^2) = 1,85$$

$$d(\mathbf{f}^1, \mathbf{f}^3) = 2,12$$



$$\mathbf{f}^2 = \begin{bmatrix} f_1^2 \\ f_2^2 \\ \vdots \\ f_{128}^2 \end{bmatrix}$$




$$\mathbf{f}^3 = \begin{bmatrix} f_1^3 \\ f_2^3 \\ \vdots \\ f_{128}^3 \end{bmatrix}$$



# CNNs Evolution

ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners





# CNNs Evolution

## Case Study: GoogLeNet

[Szegedy et al., 2014]

Deeper networks, with computational efficiency

- ILSVRC'14 classification winner (6.7% top 5 error)
- 22 layers
- Only 5 million parameters!  
12x less than AlexNet  
27x less than VGG-16
- Efficient “Inception” module
- No FC layers


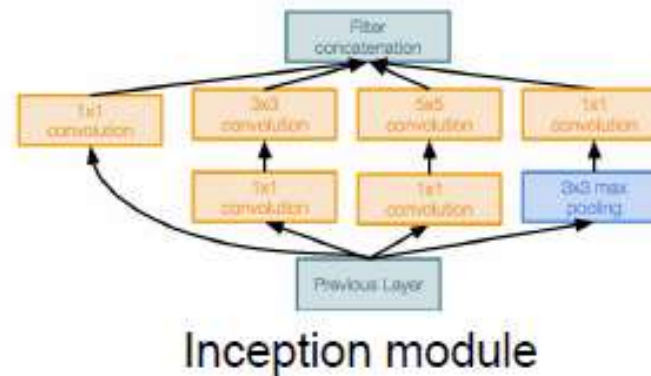


# CNNs Evolution

## Case Study: GoogLeNet

[Szegedy et al., 2014]

“Inception module”: design a good local network topology (network within a network) and then stack these modules on top of each other





# CNNs Evolution

## Case Study: ResNet

[He et al., 2015]

Very deep networks using residual connections

- 152-layer model for ImageNet
- ILSVRC'15 classification winner (3.57% top 5 error)
- Swept all classification and detection competitions in ILSVRC'15 and COCO'15!

